

# Project Writeup: Pacman Capture the Flag

Jared Jensen, Chris Gradwohl

Team: Organic Stupidity

CMPS 140 Winter 2017

University of California, Santa Cruz

## **Abstract**

We offer a q-learning offensive agent that used a handful of features exploited from the game environment. An A star search feature for an approximate q learning agent proved to be most valuable on offensive. We also provide a framework, that takes, as input, a game state instance and converts it to a network flow graph. This framework allowed us to implement a defensive agent that could discover and guard natural bottleneck positions in the game environment, making it impossible for enemies to penetrate further into defensive territory. Adding an A star attack procedure to the defensive agent further increased the agents defensive abilities.

# 1 Introduction

## 1.1 Fundamental Problem and Obstacles

Our initial approach involved understanding Capture the Flag and determining what type of agent would fit into this problem. We realized right away that the game involved several strategies based on several features. It seemed like if we could define a proper set of features then a q-learning agent would work well in Capture the Flag.

Our first goal was to create an overall offensive and defensive agent for Capture the Flag. We chose to model the Capture the Flag problem as a set of features that the q-learning agent could learn weights for. Initially we settled on *number of food pellets*, *number of defending food pellets*, *score*, *closest food pellets*, *opponent distance* as our set of features. These features seemed to be working well, and once we correctly saved the weights so that all 4 agents could work together they were performing better than expected.

## 2 Approximate Q-learning Agent

### 2.1 Agent Evaluation

In our initial implementation of the q-learning agent we implemented, did not generalize enough. We realized that when training the agent we only did so on one map. So given the feature, the agent was able to learn how to perform well given one map, but when it played on a random seed, it failed miserably.

The first problem was that we told the agent(via the feature set) to move randomly at first, and to listen to enemy movements, while it was rewarded for obtaining food pellets. After training on one map it was able to make its way out of the maze and into enemy territory to obtain food pellets and receive its reward. However given a new random map, its previous feature weights were meaningless and it was unable to get out of its initial position.

The second problem was that we implemented the weights as a global parameter, meaning that the agent would import and use the weights from a previous episode. This did not generalize to a new random map.

### 2.2 Refined Approximate Q-learning Agent

To overcome the agents mobility issues we implemented an A-star feature, that allowed the agent to learn distances to the nearest pellet. This gave the agent the ability to move into enemy territory in a more direct manner.

Secondly we removed the global weights and added a weight counter to the agent object. We did ten training games and manually added the weights the agent learned from training. We updated the agent to have a training mode so that when training it starts from scratch.

With these refinements we submitted our agent to our first tournament.

### 2.3 Notable Results

By far the most interesting(and startling) result of our q-learning agent was that given this set of features, the agent learned how to avoid ghosts. We never actually explicitly programed a feature that gave the agent a negative reward for being eaten. The feature for enemy position combined with the discount rate, allowed our agent to realize that when the *opponent distance* was really low, it performed poorly. When our agent would die it would have to go all the way to the starting part of the map, and through learning episodes

and *opponent distance* feature, it found that it was bad to be near the ghosts. In other words it obtained more rewards when it avoided ghosts, and therefore chose actions to run away! This was a really cool and interesting result of feature learning. It is a great example of choosing correct features that results in positive behavior.<sup>1</sup>

### 3 Max Flow Defense

#### 3.1 Model of the Problem

With a decent offensive q-learning agent, it seemed natural to implement a dedicated defensive agent. We needed to improve on the baseline random action reflex agent, so we considered features that could possibly benefit the defensive cause. By observation we noticed a unique bottleneck characteristic of pacman levels (see teal Ghost in Figure 1).




Figure 1: Map Bottleneck

Most levels had a spot where, the ghost could stay and completely block the enemy pacman from winning. If we could teach our agent how to find the natural bottleneck position, then that would prevent offense from consuming more pellets. Therefore, our next goal became to model the game level in such a way that allowed us to find this position given any pacman map layout.

We chose to model the map as a network flow graph, and use Ford-Fulkerson to determine the maximum flow in the graph. We constructed a framework that could convert the given into a flow network, where vertices are the possible non-wall positions and the edges are ordered pairs of possible positions.

We gave two edges (one forward and one backward) for each pair of adjacent positions. We set the capacities of all these edges to 1. We then added a source node and attached an edge with capacity  $\infty$  to each of the vertices next to the middle line<sup>2</sup>.

After several small modifications to Ford-Fulkerson and A\* we were able to find the bottleneck position and successfully find a path to its location. However it became apparent that not all bottleneck locations were created equal. That is, bottlenecks whose relative positions that were far away from the red-blue border, were not useful to guard.

Therefore we needed to find the bottleneck position with the most pacdots behind it in a particular game instance. For each pacDot, we did the following:

1. Find the max flow from the source to the pacDot, and the path through the maze that's returned by that max flow. This flow will block any paths from the source to the given dot
2. With the max flow still in place, for each vertex in the returned path, do the following:
  - (a) Try to find an augmenting path through the flow using a modified version of A\* from the source to the given vertex

<sup>1</sup>Please see the following methods on the *DummyAgent* : *update()*, *getQvalues()*, *getPolicy()*, *getFeatures()*

<sup>2</sup>Please see the *FlowNetwork* object.

- (b) If a path exists, that means there's more than one way to get to this vertex, and we continue. Otherwise, the vertex is being blocked by the max flow, and there is only one way to get to that vertex. Store this vertex.

We then end up with a list of all of the bottlenecks for all of the pacDots. We find the most commonly recurring bottleneck position, and that's the one we decide to go to. We point our defense agent to that spot and tell him to stay there as long as possible. <sup>3</sup>

## 4 Final Tweaks and Manual Override

After implementing a max flow defensive agent and correctly modeling the map as a flow network, our agent was performing well and placing in 8<sup>th</sup> in the tournaments. Our focus shifted away from the Q-learning agent and it seemed more practical in the amount of time remaining, to manually make the pacman avoid negative cases.

### 4.1 Eating Pacmen

Neither of our agents were ever told to eat pacmen, and this was a poor strategy. We programmed a procedure for both agents, to decide whether it should attack an enemy based on the A\* distance from that enemy. We noticed on offense, if an enemy was an even number of steps away from our agent, it was easy for our agent to avoid the enemy. For the defense agent, we made sure that the distance between the agent and the bottleneck is less than the distance between the enemy and the bottleneck. This allowed the agent to guard the bottleneck while still being able to attack enemies.

### 4.2 Avoiding ghosts

While our Q learning agent had figured out that it should avoid ghosts by itself, it did not avoid ghosts reliably. We integrated a parameter into our A\* search algorithm that gave a very high weight to all positions with enemy ghosts in them. Therefore the offensive agent would avoid enemy ghosts if possible.

The offensive agent also trapped itself into corners where the ghost would easily kill it. We used network flow again to determine if a ghost was nearby. We checked all nearby pacdots for one that has a maximum flow greater than 1 and therefore there is more than one path, which created less opportunities to become trapped. <sup>4</sup>

## 5 Conclusions

Although we set out to create an overall offensive/defensive agent, separating offensive and defensive concerns proved to be a valuable strategy. We realized that narrowing the problem down into smaller subproblems allowed us to focus and perfect one strategy independent of the other. The results showed that this approach is very useful when attempting to solve and model large, dynamic environments such as Capture the Flag.

---

<sup>3</sup>Please see the *findBottleneckWithMostPacdots(self, gameState)* method on the *DummyAgent* object, which calls the *FindBottlenecks(self, source, target)* method on the *FlowNetwork* object.

<sup>4</sup>Please see the *aStarSearch(self, startPosition, gameState, goalPositions, attackPacmen=True)* method on the *DummyAgent*.