

COMP90051 Statistical Machine Learning

Semester 2, 2015

Lecturer: Ben Rubinstein



THE UNIVERSITY OF
MELBOURNE

COMP90049 Revision

Covered Knowledge

- **un/supervised learning**
- **probability theory**; entropy
- association rule mining
- k -means clustering
- **naive Bayes**
- instance-based learning (IB1)
- feature selection (mutual information)
- **decision stump/tree induction** (0R, 1R, ID5)
- **basic sampling** (hold-out, cross-validation)
- **evaluation** (precision/recall/F, ROC)
- Seen: SVMs, bit of Bayes nets

Entropy

- The entropy of a discrete random event x with possible states $1, ..n$ is:


$$H(x) = - \sum_{i=1}^n P(i) \log_2 P(i)$$

where $0 \log_2 0 =^{def} 0$

Association Rules: Definitions

- An association rule is an implication $A \rightarrow B$, where A and B are disjoint itemsets
 - ★ $A = \text{antecedent}$
 - ★ $B = \text{consequent}$
- N.B. in association mining parlance:
 - ★ **item** = attribute–value pair ($I = \text{set of items}$)
 - ★ **itemset** = set of attribute–value pairs
 - ★ k -**itemset** = set of k attribute–value pairs
 - ★ **transaction** = exemplar ($T = \text{set of transactions}$)

APriori Algorithm (Rule Generation)



k -means Clustering

- Given k , the k -means algorithm is implemented in four steps:
 1. Select k points at random to act as seed clusters
 2. Compute seed points as the centroids of the clusters of the current partition (the **centroid** is the centre, i.e., mean point, of the cluster)
 3. Assign each instance to the cluster with the nearest centroid
 4. Go back to 2, stop when no reassignments
- Exclusive, deterministic, partitioning, batch clustering method

Naive Bayes (NB) Classifiers

- Classify instance $D = \langle x_1, x_2, \dots, x_n \rangle$ as class $c_j \in C$

$$c = \arg \max_{c_j \in C} P(c_j | x_1, x_2, \dots, x_n)$$

$$= \arg \max_{c_j \in C} \frac{P(x_1, x_2, \dots, x_n | c_j) P(c_j)}{P(x_1, x_2, \dots, x_n)}$$

$$= \arg \max_{c_j \in C} P(x_1, x_2, \dots, x_n | c_j) P(c_j)$$

$$= \arg \max_{c_j \in C} P(c_j) \prod_{i=1}^n P(x_i | c_j)$$

- Model trained using frequencies

Nearest Neighbour Classification

- Combining training–test instance scores to form an overall categorisation function:
- **Method 1:** index all training documents, and query the training document set with each test document; classify the test document according to the class of the top-ranked training document [**1-NN**]
- **Method 2:** index all training documents, and query the training document set with each test document; classify the test document according to the **majority class** within the k top-ranked training documents [**k-NN**]

Similarity/Distance Metrics

- Cosine similarity:

$$\text{sim}(x, y) = \frac{\vec{x} \cdot \vec{y}}{|\vec{x}| |\vec{y}|} = \frac{\sum_i x_i y_i}{\sqrt{\sum_i x_i^2} \sqrt{\sum_i y_i^2}}$$

- Relative entropy:

$$D(x \parallel y) = \sum_i x_i (\log_2 x_i - \log_2 y_i)$$

or alternatively **skew divergence**:

$$s_\alpha(x, y) = D(x \parallel \alpha y + (1 - \alpha)x)$$

Feature Selection

- **Mutual information:**

$$MI(T; C) = \sum_{t \in \{0,1\}} \sum_c P(t, c) \log_2 \frac{P(t, c)}{P(t)P(c)}$$

Constructing Decision Trees: ID3

- **Basic method:** construct decision trees in recursive divide-and-conquer fashion

FUNCTION ID3 (Root)

IF all instances at root have same class

THEN stop

ELSE Select a new attribute to use in partitioning root node instances

Create a branch for each attribute value and partition up root node instances according to each value

Call ID3($LEAF_i$) for each leaf node $LEAF_i$

- Note: we may not end up with pure leaves

Split Criteria

- The **information gain** for attribute R_A (with values x_1, \dots, x_m) at a given root node R is:

$$IG(R_A|R) = H(R) - \sum_{i=1}^m P(x_i)H(x_i)$$

- The corresponding **gain ratio** is:

$$\begin{aligned} GR(R_A|R) &= \frac{IG(R_A|R)}{H(R_A)} \\ &= \frac{H(R) - \sum_{i=1}^m P(x_i)H(x_i)}{-\sum_{i=1}^m P(x_i) \log_2 P(x_i)} \end{aligned}$$

Evaluation

- **Classification accuracy:** is the proportion of

$$\text{ACC} = \frac{TP + TN}{TP + FP + FN + TN}$$

- **Error rate:**

$$\text{ER} = \frac{FP + FN}{TP + FP + FN + TN}$$

- **Error rate reduction:**

$$\text{ERR} = \frac{\text{ER}_0 - \text{ER}}{\text{ER}_0}$$

- **Precision:**

$$\text{Precision} = \frac{TP}{TP + FP}$$

- **Recall:**

$$\text{Recall} = \frac{TP}{TP + FN}$$

- **F-score:**

$$\text{F-score} = (1 + \beta^2) \frac{PR}{R + \beta^2 P}$$

Sampling

- **Holdout** = train a classifier over a fixed training dataset, and evaluate it over a fixed held-out test dataset
- **Random Subsampling** = perform holdout over multiple iterations, randomly selecting the training and test data (maintaining a fixed size for each dataset) on each iteration
- **Cross Validation** = partition data into N folds, and use $N - 1$ as training data and 1 as test $\times N$ iterations
- **Stratified Cross Validation** = partition the data so as to maintain the overall class distribution within individual partitions

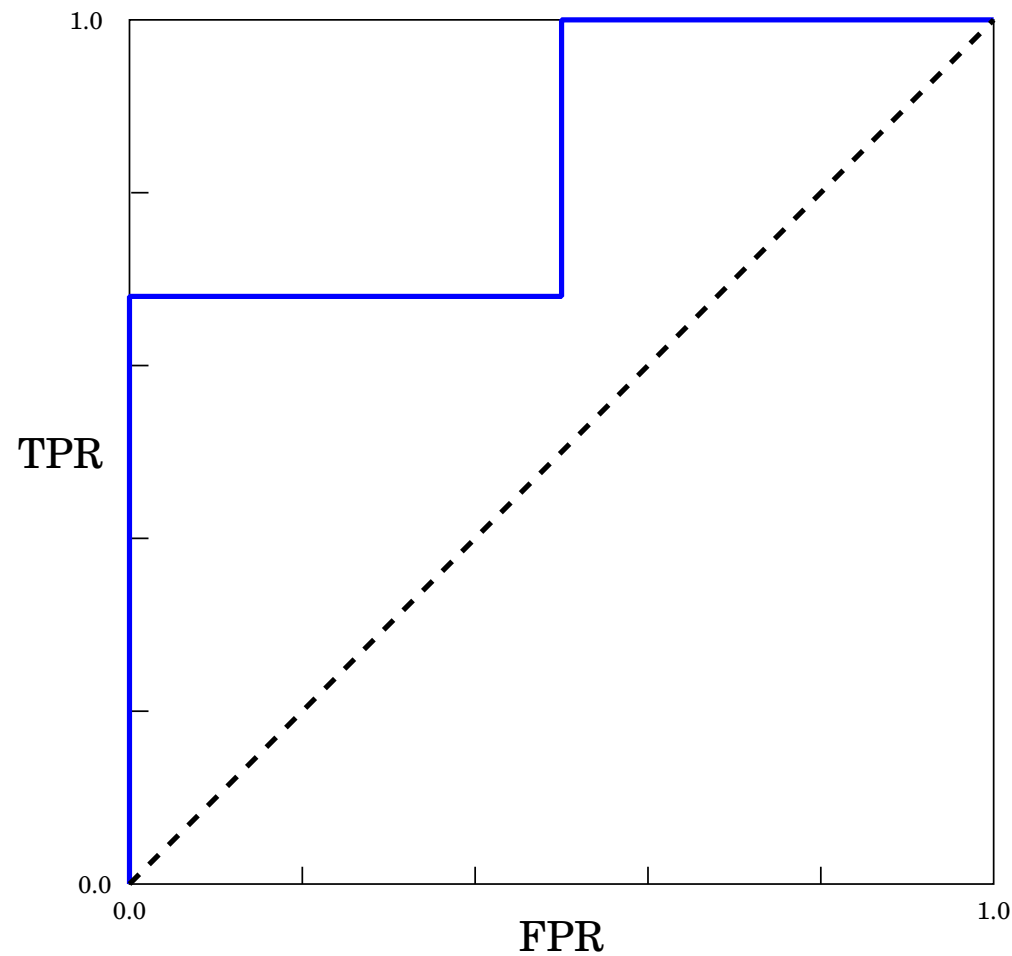
ROC Curves

1. Sort the test instances in ascending order of “rating” t_1, t_2, \dots, t_k
2. Initialise $TP_{k+1} = FP_{k+1} = 0$, and set FN_{k+1} and TN_{k+1} to the number of positive and negative instances in the dataset, resp.
3. For each $i = k, \dots, 2, 1$
 - i. update TP_i, FP_i, FN_i and TN_i assuming positive classification of instance i , based on the actual class of t_i and $TP_{i+1}, FP_{i+1}, FN_{i+1}$ and TN_{i+1}
 - ii. calculate TPR and FPR at t_i
4. Plot the TPR and FPR values for each t_i

Generating ROC Curves: Example

Class	—	+	—	+	+	
Score	0.85	0.85	0.87	0.93	0.95	
i	1	2	3	4	5	6
TP	3	3	2	2	1	0
FP	2	1	1	0	0	0
FN	0	0	1	1	2	3
TN	0	1	1	2	2	2
TPR	$\frac{3}{3}$	$\frac{3}{3}$	$\frac{2}{3}$	$\frac{2}{3}$	$\frac{1}{3}$	0
FPR	$\frac{2}{2}$	$\frac{1}{2}$	$\frac{1}{2}$	0	0	0

Generating ROC Curves: Example



**What other topics in ML have
you seen?**